



Verified Collaboration: How Lean is Transforming Mathematics, Programming, and AI

Leo de Moura
Senior Principal Applied Scientist, AWS
Chief Architect, Lean FRO

March 12, 2025



Breaking the Cycle of Uncertainty: Math, Software, and AI You Can Trust

Math, software, and AI often rely on **manual review** or **partial testing**.

An error in a theorem or critical software system can have massive consequences.

Progress dies where fear of mistakes lives.



Breaking the Cycle of Uncertainty: Math, Software, and AI You Can Trust

Math, software, and AI often rely on **manual review** or **partial testing**.

An error in a theorem or critical software system can have massive consequences.

Progress dies where fear of mistakes lives.

Lean's role: **machine-checkable proofs eliminate guesswork and create trust**.

If every step is formally verified, we unlock unprecedented confidence and collaboration.



Lean is an open-source programming language and proof assistant that is transforming how we approach mathematics, software verification, and AI.

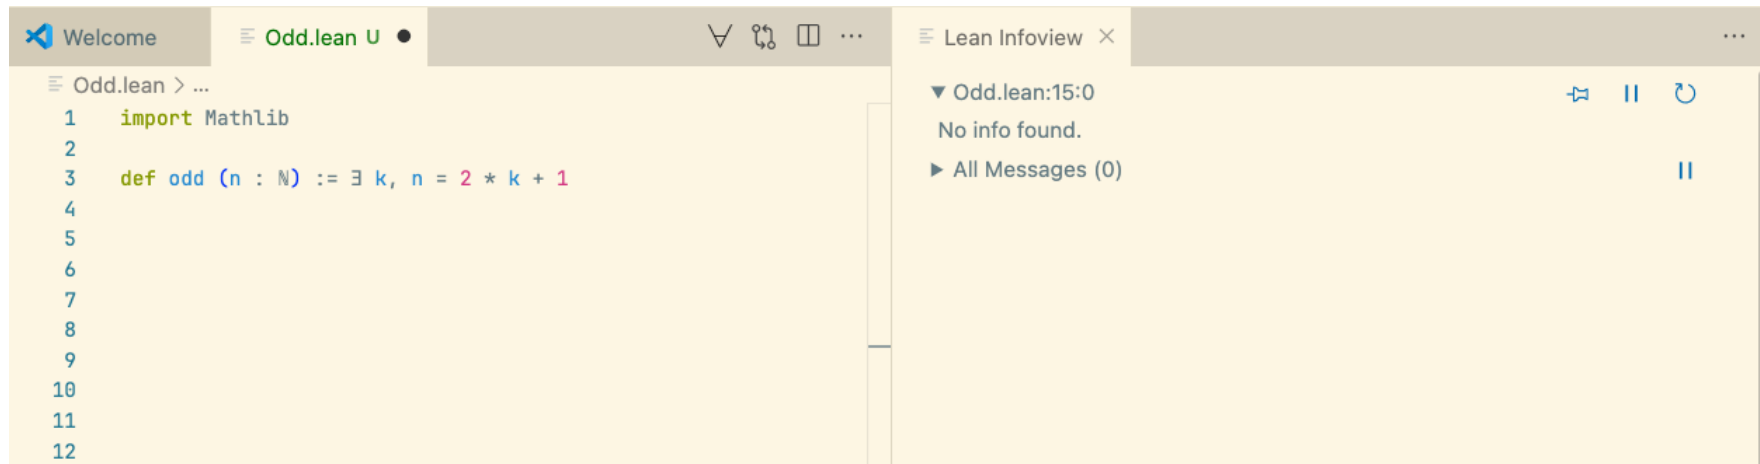
The Lean project, started in 2013, aimed at merging interactive and automated theorem proving.

Lean provides **machine-checkable proofs**.

Lean addresses the “trust bottleneck”.

Lean opens up new possibilities for collaboration.

A small example



The screenshot shows the Lean IDE interface. The top bar contains a 'Welcome' button and a tab for 'Odd.lean U'. Below the top bar, the code editor displays the following Lean code:

```
Odd.lean > ...  
1  import Mathlib  
2  
3  def odd (n : ℕ) := ∃ k, n = 2 * k + 1  
4  
5  
6  
7  
8  
9  
10  
11  
12
```

To the right of the code editor is the 'Lean Infoview' panel. It shows a message from 'Odd.lean:15:0' stating 'No info found.' Below this message is a section titled 'All Messages (0)'. The panel also includes icons for pinning, pausing, and refreshing the view.

A small example

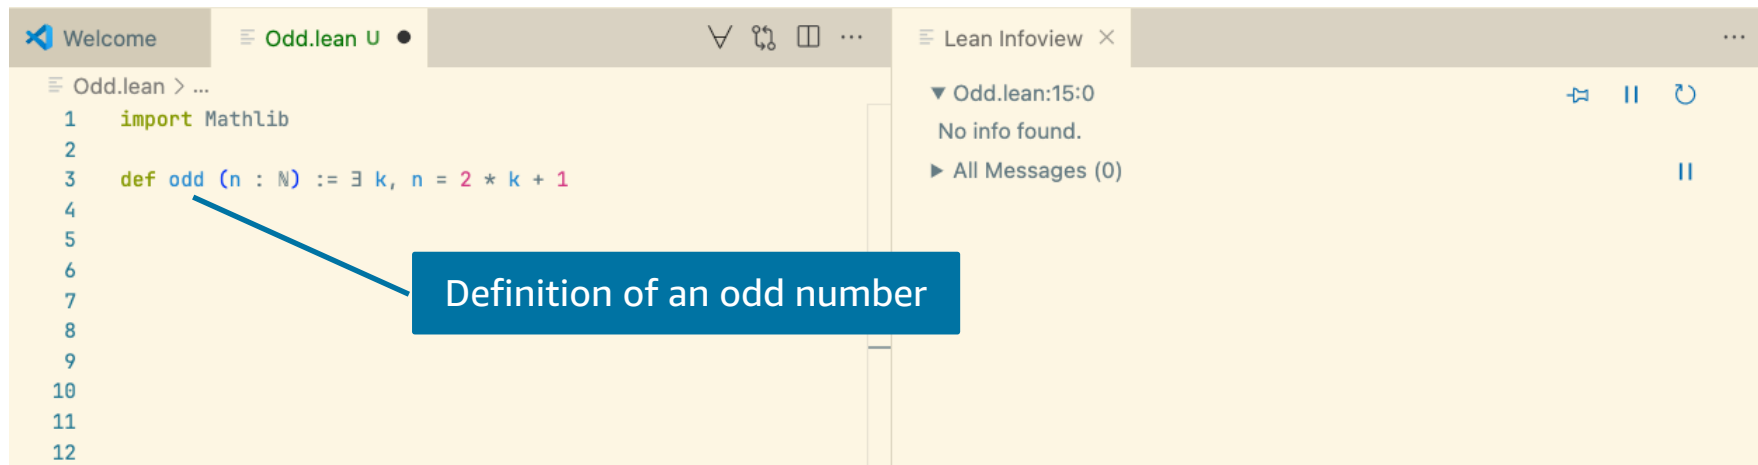
Mathlib is the Lean Mathematical library

The screenshot shows the Lean IDE interface. The top bar includes a 'Welcome' tab and a file tab 'Odd.lean U'. The main editor area displays the following code:

```
1 import Mathlib
2
3 def odd (n : ℕ) := ∃ k, n = 2 * k + 1
4
5
6
7
8
9
10
11
12
```

A blue arrow points from the 'Mathlib is the Lean Mathematical library' text box to the 'import Mathlib' line in the code. To the right, the 'Lean Infoview' panel is open, showing the message: 'Odd.lean:15:0 No info found. All Messages (0)'.

A small example



The screenshot shows the Lean IDE interface. The main editor window displays the file `Odd.lean` with the following code:

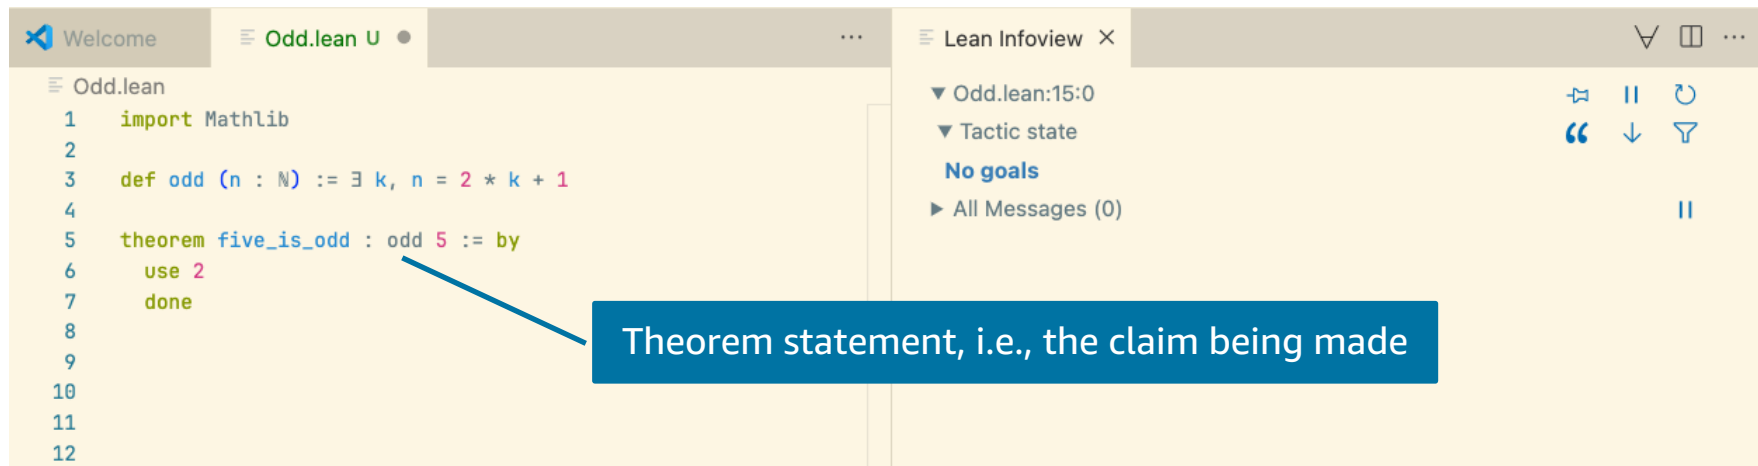
```
1 import Mathlib
2
3 def odd (n : ℕ) := ∃ k, n = 2 * k + 1
4
5
6
7
8
9
10
11
12
```

A blue callout box with the text "Definition of an odd number" points to the definition of the `odd` function on line 3.

The right-hand pane shows the "Lean Infoview" for the definition. It displays the following information:

- ▼ Odd.lean:15:0
- No info found.
- All Messages (0)

Our first theorem



The screenshot shows the Lean IDE interface. The main editor displays the file `Odd.lean` with the following code:

```
1 import Mathlib
2
3 def odd (n : ℕ) := ∃ k, n = 2 * k + 1
4
5 theorem five_is_odd : odd 5 := by
6   use 2
7   done
8
9
10
11
12
```


A blue callout box points to the theorem statement `theorem five_is_odd : odd 5 := by` with the text: "Theorem statement, i.e., the claim being made".

The right sidebar shows the `Lean Infoview` panel with the following content:

- ▼ Odd.lean:15:0
- ▼ Tactic state
- No goals**
- All Messages (0)

On the far right of the sidebar, there are icons for pinning, pausing, refreshing, quoting, scrolling, and filtering, along with a vertical scrollbar.

Our first theorem



The screenshot shows the Lean IDE interface. The main editor displays a file named `Odd.lean` with the following code:

```
1 import Mathlib
2
3 def odd (n : ℕ) := ∃ k, n = 2 * k + 1
4
5 theorem five_is_odd : odd 5 := by
6   use 2
7   done
```

A blue callout box with the text "A proof" points to the `done` keyword on line 7, indicating that the proof is complete.

The right-hand pane shows the "Lean Infoview" for the current theorem. It displays the following information:

- ▼ Odd.lean:15:0
- ▼ Tactic state
- No goals**
- All Messages (0)

On the far right of the Infoview pane, there are several icons: a pin, a pause, a refresh, a quote, a down arrow, a filter, and a double vertical bar.

Our first theorem

The screenshot shows the Lean IDE interface. On the left, the editor displays the following code in `Odd.lean`:

```
1 import Mathlib
2
3 def odd (n : ℕ) := ∃ k, n = 2 * k + 1
4
5 theorem five_is_odd : odd 5 := by
6   use 3
7   done
```

A blue callout box with the text "An incorrect proof" points to the `done` statement on line 7.

On the right, the "Lean Infoview" panel shows the tactic state:

- Odd.lean:7:2
- Tactic state
- 1 goal
- ▼ case h
- ├ 5 = 2 * 3 + 1
- Messages (1)
- All Messages (1)

Theorem proving in Lean is an interactive game

The screenshot displays the Lean IDE interface. On the left, the source code for `Odd.lean` is shown, with the following content:

```
1 import Mathlib
2
3 def odd (n : ℕ) := ∃ k, n = 2 * k + 1
4
5 -- Prove that the square of an odd number is always odd
6 theorem square_of_odd_is_odd : odd n → odd (n * n) := by
7   done
```

On the right, the `Lean Infoview` pane shows the current goal and tactic state:

```
▼ Odd.lean:7:2
▼ Tactic state
1 goal
  n : ℕ
  ⊢ odd n → odd (n * n)
► Messages (1)
► All Messages (2)
```

A blue box labeled "The 'game board'" points to the goal statement `⊢ odd n → odd (n * n)` in the tactic state.

"You have written my favorite computer game", Kevin Buzzard

Theorem proving in Lean is an interactive game

Odd.lean > `square_of_odd_is_odd`

```
1 import Mathlib
2
3 def odd (n : ℕ) := ∃ k, n = 2 * k + 1
4
5 -- Prove that the square of an odd number is always odd
6 theorem square_of_odd_is_odd : odd n → odd (n * n) := by
7   intro ⟨k1, e1⟩
8   done
9
10
11
12
```

A "game move", aka "tactic"

Lean Infoview

▼ Odd.lean:8:2

▼ Tactic state

1 goal

$n \ k_1 : \mathbb{N}$

$e_1 : n = 2 * k_1 + 1$

$\vdash \text{odd } (n * n)$

► Messages (1)

► All Messages (2)

Theorem proving in Lean is an interactive game

```
Odd.lean > square_of_odd_is_odd
1  import Mathlib
2
3  def odd (n : ℕ) := ∃ k, n = 2 * k + 1
4
5  -- Prove that the square of an odd number is always odd
6  theorem square_of_odd_is_odd : odd n → odd (n * n) := by
7    intro (k1, e1)
8    simp [e1, odd]
9    done
10
11
12
```

Lean Infoview

▼ Odd.lean:9:1

▼ Tactic state

1 goal

n k₁ : ℕ

e₁ : n = 2 * k₁ + 1


└ ∃ k, (2 * k₁ + 1) * (2 * k₁ + 1) = 2 * k + 1

► Messages (1)

► All Messages (2)

The “game move” `simp`, the simplifier, is one of the most popular moves in our game

Theorem proving in Lean is an interactive game



The screenshot displays the Lean IDE interface. On the left, the editor shows a Lean script for proving that the square of an odd number is odd. The script includes an import, a definition of 'odd', a theorem statement, and a proof using the 'by' tactic. The 'use' tactic is highlighted on line 9, with a blue arrow pointing to a callout box. On the right, the 'Lean Infoview' panel shows the current goal and tactic state. The goal is to prove the equation $(2 * k_1 + 1) * (2 * k_1 + 1) = 2 * (2 * k_1 * k_1 + 2 * k_1) + 1$. The tactic state shows the current goal and the tactics used so far.

```
Odd.lean > square_of_odd_is_odd
1  import Mathlib
2
3  def odd (n : ℕ) := ∃ k, n = 2 * k + 1
4
5  -- Prove that the square of an odd number is always odd
6  theorem square_of_odd_is_odd : odd n → odd (n * n) := by
7    intro (k₁, e₁)
8    simp [e₁, odd]
9    use 2 * k₁ * k₁ + 2 * k₁
10   done
11
12
```

Lean Infoview

Odd.lean:10:1

Tactic state

1 goal

case h

$n \ k_1 : \mathbb{N}$

$e_1 : n = 2 * k_1 + 1$

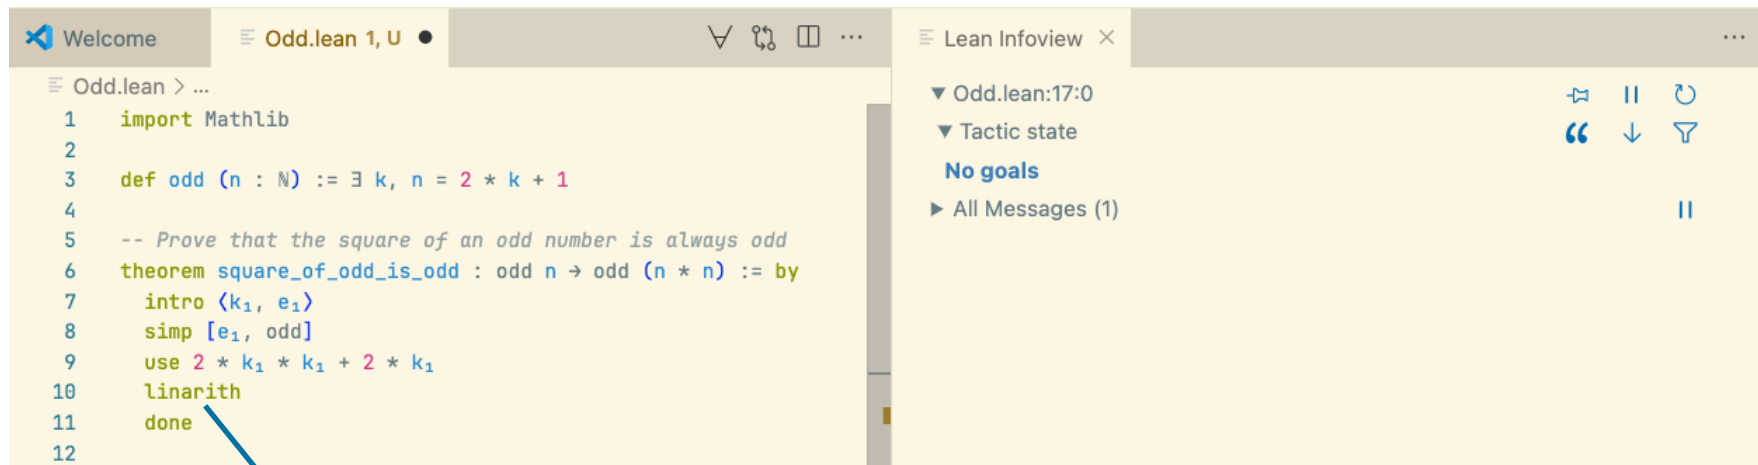
$\vdash (2 * k_1 + 1) * (2 * k_1 + 1) = 2 * (2 * k_1 * k_1 + 2 * k_1) + 1$

Messages (1)

All Messages (2)

The “game move” `use` is the standard way of proving statements about existentials

Theorem proving in Lean is an interactive game



The screenshot shows the Lean IDE interface. The main editor displays a Lean script for proving that the square of an odd number is odd. The script is as follows:

```
1 import Mathlib
2
3 def odd (n : ℕ) := ∃ k, n = 2 * k + 1
4
5 -- Prove that the square of an odd number is always odd
6 theorem square_of_odd_is_odd : odd n → odd (n * n) := by
7   intro (k₁, e₁)
8   simp [e₁, odd]
9   use 2 * k₁ * k₁ + 2 * k₁
10  linarith
11 done
12
```


The right-hand pane shows the 'Lean Infoview' with the following content:

- Odd.lean:17:0
- Tactic state
- No goals
- All Messages (1)

Navigation icons are visible on the right side of the Infoview pane.

We complete this level using `linarith`, the linear arithmetic, move

Theorem proving in Lean is an interactive **and addictive** game



The screenshot shows the Lean IDE interface. The left pane displays the source code for a file named `Odd.lean`. The code defines an odd number and proves that its square is also odd. The right pane shows the Lean Infoview, which indicates that there are no goals left to prove.

```
1 import Mathlib
2
3 def odd (n : ℕ) := ∃ k, n = 2 * k + 1
4
5 -- Prove that the square of an odd number is always odd
6 theorem square_of_odd_is_odd : odd n → odd (n * n) := by
7   intro (k₁, e₁)
8   simp [e₁, odd]
9   use 2 * k₁ * k₁ + 2 * k₁
10  linarith
11  done
12
```

Lean Infoview: Odd.lean:17:0
Tactic state
No goals
All Messages (1)

*"You can do 14 hours a day in it and not get tired and feel kind of high the whole day.
You're constantly getting positive reinforcement", Amelia Livingston*

Mathlib

The Lean Mathematical Library supports a wide range of projects.

It is an open-source **collaborative project** with over 500 contributors and 1.7M LoC.

"I'm investing time now so that somebody in the future can have that amazing experience",

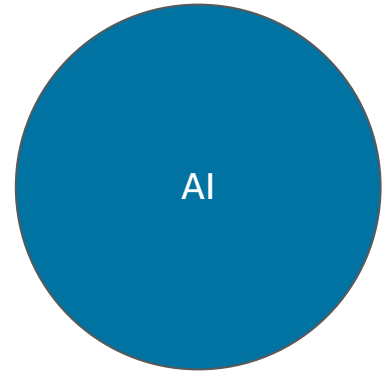
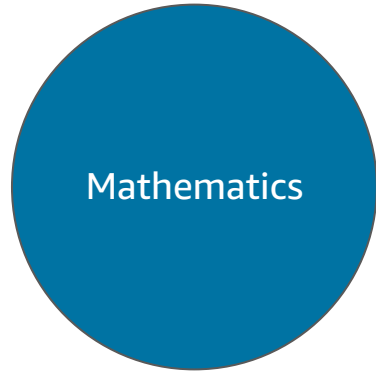
Heather Macbeth

KEVIN HARTNETT SCIENCE OCT 11, 2020 8:00 AM

The Effort to Build the Mathematical Library of the Future

A community of mathematicians is using software called Lean to build a new digital repository. They hope it represents where their field is headed next.





Mathematics

Preamble: the Perfectoid Spaces Project

Kevin Buzzard, Patrick Massot, Johan Commelin

Goal: Demonstrate that we can **define complex mathematical objects** in Lean.


They translated Peter Scholze's definition into a form a computer can understand.

It not only achieved its goals but also demonstrated to the math community that **formal objects can be visualized and inspected with computer assistance**.

Math is now **data** that can be **processed, transformed**, and **inspected** in various ways.

Preamble: the Perfectoid Spaces Project (cont.)

Kevin Buzzard, Patrick Massot, Johan Commelin



mathoverflow

Home

What are "perfectoid spaces"?



Here is a completely different kind of answer to this question.

72

A *perfectoid space* is a term of type `PerfectoidSpace` in the [Lean theorem prover](#).



Here's a quote from the source code:

```
structure perfectoid_ring (R : Type) [Huber_ring R] extends Tate_ring R : Prop :=
  (complete : is_complete_hausdorff R)
  (uniform : is_uniform R)
  (ramified : ∃ w : pseudo_uniformizer R, w^p ∉ p in R)
  (Frobenius : surjective (Frob R/p))
```

The Challenge

In November of 2020, Peter Scholze posits the Liquid Tensor Experiment (LTE) challenge.

*"I spent much of 2019 **obsessed** with the proof of this theorem, **almost getting crazy over it**. In the end, we were able to get an argument pinned down on paper, but I think nobody else has dared to look at the details of this, and so I still have some small lingering doubts",*

Peter Scholze

The First Victory

Johan Commelin led a team with several members of the **Lean community** and announced the **formalization of the crucial intermediate lemma** that Scholze was unsure about, with only minor corrections, in **May 2021**.

"[T]his was precisely the kind of oversight I was worried about when I asked for the formal verification. [...] The proof walks a fine line, so if some argument needs constants that are quite a bit different from what I claimed, it might have collapsed", Peter Scholze

nature

[Explore content](#) [Journal information](#) [Publish with us](#) [Subscribe](#)

[nature](#) > [news](#) > [article](#)

NEWS | 18 June 2021

**Mathematicians welcome
computer-assisted proof in ‘grand
unification’ theory**

Achieving the Unthinkable

The full challenge was completed in July 2022.

The team not only verified the proof but also simplified it.

Moreover, they did this without fully understanding the entire proof.

Johan, the project lead, reported that he could only see two steps ahead. **Lean was a guide.**

“The Lean Proof Assistant was really that: an assistant in navigating through the thick jungle that this proof is. Really, one key problem I had when I was trying to find this proof was that I was essentially unable to keep all the objects in my RAM, and I think the same problem occurs when trying to read the proof”, Peter Scholze

Only the Beginning

Independence of the Continuum Hypothesis, Han and van Doorn, 2021

Sphere Eversion, Massot, Nash, and van Doorn, 2020-2022

Fermat's Last Theorem for regular primes, Brasca et al., 2021-2023

Unit Fractions, Bloom and Mehta, 2022

Consistency of Quine's New Foundations, Wilshaw and Dillies, 2022-2024

Polynomial Freiman-Ruzsa Conjecture (PFR), Tao and Dillies, 2023

Prime Number Theorem And Beyond, Kontorovich and Tao, 2024-ongoing

Carleson Project, van Doorn, 2024-ongoing

Fermat's Last Theorem (FLT), Buzzard, 2024-ongoing, community estimates it will take +1M LoC

What did we learn?

Machine-checkable proofs enable a new level of **collaboration** in mathematics.

The power of the **community**.

It is not just about proving but also understanding complex objects and proofs, getting new insights, and navigating through the “thick jungles” that are **beyond our cognitive abilities**.

Software

Lean in Software Verification: The Story of SampCert

Lean is a programming language, and is used in **many software verification projects**.

You can write code and reason about it simultaneously.

You can prove that your code has the properties you expect.

“Testing can show the presence of bugs, but not their absence”, E. Dijkstra

Differential Privacy

A mathematical framework that ensures the **privacy of individuals** in a dataset by adding controlled **random noise** to the data.

Discrete sampling algorithms, like the **Discrete Gaussian Sampler**, are used to add carefully calibrated noise to data.

What may go wrong if a buggy sampler is used?

Privacy Violations: leakage of sensitive information

Incorrect Results: distorted analysis results



SampCert

A project led by **Jean-Baptiste Tristan** at AWS.

An **open-source** Lean library of formally **verified differential privacy primitives**.

Tristan's implementation is not only verified, but it is also **twice as fast as the previous one**.

He managed to implement **aggressive optimizations** because Lean served as a guide, ensuring that **no bugs** were introduced.

AWS Clean Rooms Differential Privacy

Protect the privacy of your users with mathematically backed controls in a few steps

SampCert would not exist without Mathlib

SampCert is software, but its verification relies heavily on Mathlib.

The verification of code addressing practical problems in data privacy depends on the formalization of mathematical concepts, from **Fourier analysis** to **number theory** and **topology**.

Many more open-source projects

Cedar, a policy language and evaluation engine.

LNSym, a symbolic simulator for Armv8 native-code programs: cryptographic machine-code programs.

TenCert, a tensor compiler, verified StableHLO and NKI.

NFA2SAT, a verified string solver.

Many more at the **Lean Project Registry**: <https://reservoir.lean-lang.org/>



What did we learn?

Machine-checkable proofs enable you to **code without fear**.

AI



Lean Enables **Verified** AI for Mathematics and Code

LLMs are powerful tools, but they are prone to **hallucinations**.

In Math, a **small mistake can invalidate the whole proof**.

Imagine manually checking an AI-generated proof with the size and complexity of FLT.

The informal proof is **over 200 pages**.

Buzzard estimates a formal proof will require more than **1M LoC** on top of Mathlib.

Machine-checkable proofs are the antidote to hallucinations.

AI Proof Assistants


Several groups are developing AI that suggests the **next move**(s) in Lean's interactive proof game.

[LeanDojo](#) is an open-source project from Caltech, and everything (model, datasets, code) is open.

[OpenAI](#) and [Meta AI](#) have also developed AI assistants for Lean.

AI Proof Assistants

LeanCopilot is part of the LeanDojo project at Caltech. It uses the move (aka tactic) suggestion feature available in the Lean IDE.



The screenshot displays the Lean IDE interface. The main editor shows a Lean 4 file named `Lean4Example.lean` at line 5, column 1. The code defines an example `example (a b c : Nat) : a * (b + c) = a * c + a * b := by` and then uses the `suggest_tactics` tactic, which is underlined in red. The `done` tactic is also underlined in red. The Lean Infoview panel on the right shows the current goal state and suggests tactics for proving the goal.

```
1 import LeanCopilot
2
3 example (a b c : Nat) : a * (b + c) = a * c + a * b := by
4   suggest_tactics
5   done
```

Lean Infoview

▼ Lean4Example.lean:4:2

▼ Tactic state

1 goal

a b c : Nat

⊢ a * (b + c) = a * c + a * b

▼ Tactic suggestions

Try this:

- `simp [Nat.left_distrib, Nat.add_comm]`
- `rw [Nat.mul_add, Nat.add_comm]`
- `rw [Nat.mul_add, Nat.mul_comm, Nat.add_comm]`
- `simp [Nat.add_comm]`
- `rw [Nat.mul_comm, Nat.add_comm]`
- `rw [Nat.mul_comm, Nat.mul_comm]`
- `rw [Nat.mul_add, Nat.mul_comm]`
- `apply Nat.add_left_cancel`

Move Over, Mathematicians, Here Comes AlphaProof

A.I. is getting good at math — and might soon make a worthy collaborator for humans.

Share full article 47




Ring the gong at Google Deepmind's London headquarters, a ritual to celebrate each A.I. milestone, including its recent triumph of reasoning at the International Mathematical Olympiad. Google Deepmind

Auto-formalization

The process of converting natural language into a formal language like Lean.

It is much **easier to learn to read Lean than to write it**.

[LeanAide](#) is one of the auto-formalization tools available for Lean.



```
LeanTimes.lean PnP2023/Extras/LeanTimes.lean
You, 5 days ago | 1 author (You)
import Mathlib
import LeanAide
💡
/- There are infinitely many odd numbers -/
/- Every prime number is either `2` or odd -/

Lean Infoview
▼ LeanTimes.lean:4:43
No info found.
► All Messages (0)
```

What did we learn?

Machine-checkable proofs enable **AI that does not hallucinate**.

LLMs enable **auto-formalization**.

Machine learning opens doors to **new proof search engines**.

Before we wrap up...



Lean Enables Decentralized Collaboration

Lean is Extensible

Users extend Lean using Lean itself.

Lean is implemented in Lean.

You can make it your own.

You can create your own moves.


Machine-Checkable Proofs

You don't need to trust me to use my proofs.

You don't need to trust my automation to use it.

Code without fear.

Lean is a game where we can implement your own moves



The screenshot shows the Lean IDE interface. On the left, a code editor displays a proof script in `Odd.lean`. The script defines an `odd` function and proves a theorem about the square of an odd number. The `linarith` tactic is used to complete the proof. On the right, the 'Lean Infoview' panel shows the current tactic state, which is empty, indicating the proof is complete.

```

1  import Mathlib
2
3  def odd (n : ℕ) := ∃ k, n = 2 * k + 1
4
5  -- Prove that the square of an odd number is always odd
6  theorem square_of_odd_is_odd : odd n → odd (n * n) := by
7    intro ⟨k₁, e₁⟩
8    simp [e₁, odd]
9    use 2 * k₁ * k₁ + 2 * k₁
10   linarith
11   done
12

```

The right panel shows the tactic state:


- Odd.lean:17:0
- Tactic state
- No goals
- All Messages (1)

The `linarith` “move” was implemented by the Mathlib community in Lean!

You can use Lean to introspect its internal data

The tool [lean-training-data](#) is implemented in Lean itself. **It is a Lean package.**

A similar approach can be used to automatically generate proof animations.





Lean FRO: Shaping the Future of Lean Development

The Lean Focused Research Organization (FRO) is a non-profit dedicated to Lean's development.

Founded in **August 2023**, the organization has 18 members.

Its mission is to enhance critical areas: **scalability, usability, documentation**, and **proof automation**.



It must reach **self-sustainability in August 2028** and become the **Lean Foundation**.

Philanthropic support is gratefully acknowledged from the **Simons Foundation**, the **Alfred P. Sloan Foundation**, **Richard Merkin**, and **Founders Pledge**.

FROs accelerate scientific progress

James Webb Telescope and CERN illustrate a common pattern in science: a need for projects that are bigger than an academic lab can undertake, more coordinated than a loose consortium or themed department, and not directly profitable enough to be a venture-backed startup or industrial R&D project.


<https://www.convergentresearch.org/about-fros>



Growth of Lean projects on GitHub



New installations of Lean Development Environment (2024 to present)





How can I contribute?

Help building [Mathlib](#).

Want to engage with the vibrant Lean community? Join our [Zulip channel](#).

Interested in ML kernels? Contribute to the [KLR project](#).

Want to contribute to a large formalization project? Join the [FLT formalization project](#).

Start your own open-source Lean project! Your package will be available on our registry [Reservoir](#).

Conclusion


Lean is an **efficient programming language** and **proof assistant**.

The Mathlib community is changing how math is done.

It is not just about proving but also understanding complex objects and proofs, getting new insights, and navigating through the “thick jungles” that are **beyond our cognitive abilities**.

Lean tracks details, so humans focus on big ideas.

Decentralized collaboration with Lean: Large teams can collectively tackle huge proofs without losing track.



Thank You

<https://leanprover.zulipchat.com/>

x: @leanprover

LinkedIn: Lean FRO

Mastodon: @leanprover@functional.cafe

#leanlang, #leanprover

<https://www.lean-lang.org/>

